

L05 How a neural network learns

A network is a function with many unknown numbers, fitted to examples the way you would fit a line.

LiDAR scans

finishing Monday: depth sensors

- **Spinning** units sweep 360° with a fixed set of elevation beams. **Solid-state** units have no moving parts, a narrower field of view, and a lower price.
- Each beam traces a ring, so the cloud is dense along a ring and sparse between rings: a 32-beam unit at 20 m has rings about 40 cm apart.
- Points arrive over the sweep, so a moving robot measures one scan from many poses.





Image: Shvicha, CC BY-SA 4.0, via Wikimedia Commons

Metric versus relative depth

Metric depth is in metres. Relative depth gets the ordering and the shape right but not the size, and a planner or a gripper needs the size.

Source	What you get	Where the metres come from
Stereo	metric	the baseline B , measured once
LiDAR, time of flight	metric	the speed of light and a time
Monocular SfM or SLAM	metric up to one global scale	nowhere: no image fixes it
Monocular depth network	relative: scale and shift unknown per image	nowhere: trained to ignore them
Metric depth network	approximately metric	the intrinsics, or training on one camera

Where learned depth fits

- **Monocular depth networks** such as MiDaS and Depth Anything predict depth from one image, and they are good at *relative* depth.
- Most are trained with a scale- and shift-invariant loss, so the output is correct only up to $aZ + b$, with a and b unknown for each image.
- **Metric** variants take the scale from the intrinsics or from training on one camera; **depth completion** takes it from sparse LiDAR returns. Either way it comes from the sensor or the body, never from the image.





Figure: MiDaS (Intel ISL), MIT License



How a neural network learns

L05 · Unit 0

A depth network is a function with millions of unknown numbers, found from examples the way you would fit a line.

Nothing today assumes you have trained a network before. It starts from $\hat{y} = wx + b$, works one training step by hand, and ends with the encoder–decoder shape of the depth networks on the last slide.

Today

1. **One neuron:** a line with a weight and a bias.
2. **Learning:** guess, check, improve. A loss, gradient descent, and backpropagation, worked by hand.
3. **Layers and activations:** why a stack of lines needs a bend.
4. **Classification:** softmax and cross-entropy.
5. **Training in practice:** batches, epochs, validation, overfitting, and keeping training stable.
6. **Convolutional networks:** filters, stride, padding, pooling, and the shape of a depth network.

Learning a function from examples


- A **neural network** is a function $\hat{y} = f_{\theta}(x)$ with many unknown **parameters** θ : its weights and biases.
- **Supervised learning**: given pairs (x, y) of inputs and correct answers, choose θ so that $f_{\theta}(x)$ lands near y .
- SIFT and ORB (L04) were functions designed by hand. A network's numbers are found from data.



Image: Josef Steppan, CC BY-SA 4.0, via Wikimedia Commons

You choose the form of the function and the loss that scores a wrong answer; training on data picks the numbers.

One neuron is a line



$$\hat{y} = wx + b$$


- w is the **weight** and b the **bias**.
- The **forward pass** computes $\hat{y}$ from x with the current w and b .
- Rosenblatt's 1958 **perceptron** was a unit like this, followed by a threshold.



Photo: U.S. Navy, public domain, via Wikimedia Commons

If you knew w and b you would be done. Training is how to find them when you don't.

Many inputs: a weighted sum



$$\hat{y} = w_1x_1 + w_2x_2 + w_3x_3 + b = \mathbf{w}^T \mathbf{x} + b$$

- One weight per input, one bias per neuron.
- The name comes from biology: a neuron sums the signals arriving at its dendrites and fires. The resemblance stops about there.






Figure: Stanford CS231n notes, MIT License

Guess, check, improve



- **Training data:** the (x, y) pairs used to fit the weights.
- **Loss:** one number that says how wrong the predictions are.
- **Backpropagation:** work out how much each weight contributed to the loss, then change every weight to reduce it.
- Start from small random weights, drawn from a distribution. A guess is fine; the loop fixes it.


The loss measures the miss



$$L(w, b) = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2 = \frac{1}{N} \sum_{i=1}^N (wx_i + b - y_i)^2$$

- This is the **mean squared error**: squaring makes a large miss cost far more than a small one.
- The data are fixed, so L is a function of w and b alone. Training means making it small.

Gradient descent: walk downhill



$$w \leftarrow w - \eta \frac{\partial L}{\partial w}, \quad b \leftarrow b - \eta \frac{\partial L}{\partial b}$$





Figure: Google Machine Learning Crash Course, CC BY

4.0

- Positive slope: decrease w . Negative slope: increase it. The **learning rate** η sets the step size.
- Over w and b the loss is a surface (right). With a million weights, each still moves against its own slope.

The learning rate




- **Too small:** every step is safe and training takes forever.
- **Too large:** the steps overshoot the bottom and bounce, and noisy data makes it worse. Larger still and the loss diverges.
- Tune it before any other setting. Starting points between 10^{-1} and 10^{-4} are common.

One training step, by hand

1 of 2: forward pass and gradients

One example: $x = 2.1$, target $y = 4$. First guess: $w = 1$, $b = 0$.


■ The loss, then both gradients by the chain rule



One training step, by hand


2 of 2: update, and check

■ Take one step with $\eta = 0.01$, then recompute the loss



$\eta = 0.01$ is large for a real network and fine for a toy problem with one example.

Backpropagation is the chain rule on a graph



- **Forward:** compute every node from its inputs, left to right, and keep the values.
- **Backward:** start at the loss and multiply by each node's local derivative on the way back. A node only needs the values stored at its own inputs.
- One backward pass gives the gradient for every weight at once. PyTorch builds this graph as your code runs, and `loss.backward()` walks it.

Stacking lines gives a line

Two neurons in a row, then two side by side, with nothing between them





Image: Naokin0525, CC0, via Wikimedia Commons

Any stack of linear layers collapses to a single linear layer, and a linear classifier can only split the data with a straight line. No line separates XOR.

An activation function bends the line




$$\hat{y} = \sigma(w^\top x + b)$$

Apply a nonlinear function σ to every neuron's sum, and layers no longer collapse.




one unit

$$\max(0, wx + b)$$



5 units

each orange dot is one bend




14 units

dashed: the function to fit

Each ReLU unit adds one bend. With enough units a single hidden layer can follow almost any continuous function: networks are **universal approximators**.


Three common activations




	Formula	Range	Typical use
ReLU	$\max(0, x)$	$[0, \infty)$	hidden layers: the default
tanh	$(e^x - e^{-x}) / (e^x + e^{-x})$	$(-1, 1)$	hidden layers in older networks
sigmoid	$1 / (1 + e^{-x})$	$(0, 1)$	the output for a yes-or-no probability

Choosing an activation

- **Range:** bounded or not? A depth output should not be squashed into $(0, 1)$.
- **Differentiability:** backpropagation needs a derivative. ReLU's kink at zero does no harm.
- **Cost:** ReLU is a threshold, and it zeroes negative sums, so fewer units are active.
- **Gradient flow:** the sigmoid's slope is at most 0.25; ReLU's is 1 wherever it is active.
- **The problem:** sigmoid for yes or no, softmax for one of many classes, none for regression.



A multilayer network



$$h_1 = \sigma(W_1x + b_1)$$

$$h_2 = \sigma(W_2h_1 + b_2)$$


$$\hat{y} = W_3h_2 + b_3$$

- Each layer is a weight matrix, a bias vector and an activation.
- **Fully connected:** every neuron reads every output of the layer before.
- **Hidden** layers are the ones between input and output. **Deep** means many of them.

Counting parameters

A digit classifier: the 28×28 image flattened to 784 inputs, two hidden layers of 16, and 10 outputs, one per digit.

■ Count the weights and biases of a 784–16–16–10 network




That is the network in 3Blue1Brown's series. Depth Anything's large model has about 335 million.

Watch one learn


playground.tensorflow.org

1. Pick the spiral dataset. Set the activation to **linear** and train: the boundary stays straight.
2. Switch to **ReLU**. Add a hidden layer and neurons until the spiral separates.
3. Raise the learning rate to 3 and watch the loss jump.



Screenshot: TensorFlow Playground, Apache-2.0

Classification: scores to probabilities



$$p_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$$

- The last layer gives one **score** z_k per class. Scores can be any real number.
- **Softmax** makes them positive and makes them sum to one, and it keeps their order.

Cross-entropy loss

$$L = -\log p_{\text{correct}}$$

■ The loss for the scores on the last slide

- Confident and right costs almost nothing.
Confident and wrong costs a lot.
- Squared error for quantities such as depth;
cross-entropy for categories.




The output layer matches the task

Task	Robot example	Output layer	Loss
Regression	depth at every pixel	none, or ReLU to keep it positive	squared or absolute error
Yes or no	is this pixel free space?	sigmoid	binary cross-entropy
One of many	which of 20 object classes?	softmax	cross-entropy

Semantic segmentation is classification at every pixel, and a depth network is regression at every pixel. The hidden layers can be identical.

Batches and epochs



- A gradient over the whole dataset for every step is too slow. Average it over a random **mini-batch** instead: **stochastic gradient descent**.
- **Epoch**: one complete pass of the training data through the network. 60,000 MNIST images in batches of 100 is 600 steps per epoch.
- Batch size is capped by GPU memory, and 3D data fills memory quickly.

The training loop, in PyTorch

```
model = nn.Sequential(nn.Flatten(),
                      nn.Linear(784, 16), nn.ReLU(),
                      nn.Linear(16, 16), nn.ReLU(),
                      nn.Linear(16, 10))

loss_fn = nn.CrossEntropyLoss()      # softmax, then cross-entropy
opt = torch.optim.SGD(model.parameters(), lr=0.1)

for epoch in range(10):
    for x, y in train_loader:         # one mini-batch at a time
        loss = loss_fn(model(x), y)   # check: forward pass and loss
        opt.zero_grad()               # clear the last step's gradients
        loss.backward()                # backpropagation
        opt.step()                     # improve:  $w = w - lr * \text{gradient}$ 
```

Guess is the random initialisation inside `nn.Linear`. The inner loop is check and improve, once per mini-batch.

Did it learn? Three datasets



- **Training:** used to update the weights.
- **Validation:** never used to update weights, but evaluated periodically during training, to choose the learning rate, the model size, and when to stop.
- **Test:** evaluated once, after every choice is made. It estimates how the model does on data it has never seen.

Trap


Split robot data by scene or by drive, never by frame. Neighbouring frames are nearly the same image, so a random split puts the test set inside the training set.

Reading the loss curves




1. Training loss falls, validation loss does not move: the model is fitting the training set and learning nothing that carries over.
2. Both fall: it is learning something about what you are trying to model. Keep going.
3. Training falls while validation rises: **overfitting**. Stop, and keep the checkpoint with the lowest validation loss (**early stopping**).


Overfitting



underfit



about right



overfit





Image: Chabacano, CC BY-SA 4.0,
via Wikimedia Commons

- A model with enough parameters can pass through every training point, noise included, and be wrong everywhere else.
- **Fixes:** more data; **data augmentation** (flips, crops, colour changes); a smaller model; **dropout**; **weight decay**, a penalty on large weights; early stopping.

Vanishing and exploding gradients



- The gradient reaching an early layer is a product of one factor for every layer after it.
- Factors below 1 shrink it: **vanishing gradients**, and early layers stop learning. With sigmoid slopes of at most 0.25, ten layers give $0.25^{10} \approx 10^{-6}$.
- Factors above 1 blow it up: **exploding gradients**, often helped along by unbounded activations, and the loss jumps to NaN.
- **Dead ReLU**: a unit whose sum is negative for every input outputs 0 with slope 0, and never updates again.


Keeping training stable

- **ReLU or a variant** in hidden layers. Leaky ReLU keeps a small slope below zero, so its units cannot die.
- **A lower learning rate:** large steps are what push ReLU units dead and gradients up.
- **Sensible starting weights**, random and scaled to the number of inputs of each layer.
- **Normalisation** of the input data, and of the outputs of intermediate layers.


Activation, learning rate and normalisation are chosen together. Change one and the best setting of the others moves.

Normalisation

■ Batch normalisation of one feature over a batch of B examples



one feature,
over the batch



one example,
over its features

rows: examples in a batch • columns: features

Normalisation rescales numbers to control their mean and spread. For input images: subtract the dataset mean and divide by its standard deviation, per colour channel.

Batch norm or layer norm


- **Batch norm** averages each feature over the batch. At test time there may be one image, so it uses running averages saved during training. Forgetting `model.eval()` is a classic bug.
- It degrades at small batch sizes, and 3D perception often trains with batches of 2 to 8.
- **Layer norm** averages over the features of one example, so it behaves the same at any batch size, including 1. Transformers use it.

With batches of a few scenes, reach for layer norm or group norm.

An image in a fully connected layer


■ Count the weights: a 224×224 colour image into 1000 neurons

one pixel: a brightness



- Edges, corners and textures live in patches of pixels. A single pixel is only a brightness.
- A fully connected layer has no idea which pixels are neighbours: shuffle every image's pixels the same way and it learns just as well.

Convolution: slide a small filter



$$(I * K)(r, c) = \sum_m \sum_n I(r + m, c + n) K(m, n)$$

- Lay the filter on a patch, multiply entry by entry, add. The shaded output is $(3 + 2 + 3) - (0 + 1 + 0) = 7$: the patch's left column minus its right. Slide one pixel and repeat.
- The same nine weights are used at every position: **weight sharing**. An edge detector works anywhere in the image.

Filters you have already met

1	2	1
2	4	2
1	2	1

blur
× 1/16

-1	0	1
-2	0	2
-1	0	1

vertical edges
Sobel

w_{11}	w_{12}	w_{13}
w_{21}	w_{22}	w_{23}
w_{31}	w_{32}	w_{33}

learned
set by backprop






Figure: Olah et al., Distill (2020), CC BY 4.0

- L04's image pyramid blurred with a Gaussian, $G(x, y) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$, sampled on a small grid. A bigger filter or a larger σ blurs more.
- An edge filter responds where brightness changes from one side to the other.
- SIFT and ORB were built from filters like these, designed by hand. A convolutional network treats the nine numbers as weights and learns them. On the right, the first-layer filters of a trained network: edges and colour blobs.


Stride and padding



stride 1
output 5 wide



stride 2
output 3 wide




padding 1
5 wide in, 5 wide out

- **Stride:** how far the filter moves each time. Larger strides give smaller outputs; above 2 is rare.
- **Padding:** a border of zeros around the input, to control the output size.

■ Output width for input W , filter K , padding P , stride S


Channels: many filters at once



- A colour image has 3 **channels**. A filter spans all of them: $3 \times 3 \times 3$ weights and a bias.
- Use C_{out} filters and the output has C_{out} channels, one **feature map** per filter. This is the layer's **depth**.


■ Parameters: 3×3 filters, 3 channels in, 64 out

Receptive field




- The **receptive field** of a neuron is the patch of the input it depends on.
- Each 3×3 layer widens it by 2, so three layers see 7×7 . Stride and pooling widen it faster.
- How large it must be depends on the task. Fine geometry needs small neighbourhoods; judging motion or the layout of a room needs large ones.

Pooling




- **Max pooling** keeps the largest value in each window; **average pooling** keeps the mean.
- A 2×2 window with stride 2 halves height and width, and has no weights to learn.
- It downsamples the feature maps between layers: less computation later, and a wider receptive field.

Dropout



- During training, zero each neuron's output with probability p , often 0.5, and draw a new random set at every step.
- No neuron can rely on one particular partner being there, which reduces overfitting.
- At test time every neuron is used, with outputs scaled to match what the next layer saw in training.

A convolutional network



- Convolutions and pooling shrink height and width and add channels: more kinds of feature, each summarising a wider area.
- Then flatten, finish with fully connected layers, and softmax over the classes.
- This is LeNet (LeCun et al., 1998), which read handwritten digits. AlexNet (2012) used the same recipe, far larger, on GPUs, and won ImageNet.


What the layers learn



Figure: Olah, Mordvintsev & Schubert, Distill (2017), CC BY 4.0

- Early filters respond to edges and colour; later layers to textures, patterns, parts, and whole objects.
- A trained CNN has learned its own image pyramid and its own feature descriptors, tuned to its task.

From a classifier to a depth network



Drop the fully connected layers. An **encoder** shrinks the image into features, a **decoder** grows them back to one output per pixel, and **skip connections** carry fine detail across (U-Net). Train it with a loss at every pixel against measured depth.

The loss decides what a network can learn. MiDaS's loss ignores scale and shift, so its depth is relative.

What you should be able to do

- Run one gradient-descent step on $\hat{y} = wx + b$ by hand, chain rule included.
- Say why a network needs a nonlinear activation, and pick one for the hidden layers and one for the output.
- Turn scores into probabilities with softmax, and score them with cross-entropy.
- Read a pair of loss curves and decide whether to keep training.
- Compute the output size and the parameter count of a convolutional layer.

To go further: 3Blue1Brown, *Neural networks*, chapters 1–4; *Dive into Deep Learning* (d2l.ai), chapters 3–5 and 7; the Stanford CS231n notes on convolutional networks.

Unit 0 is done

Unit 1 begins next lecture

Five lectures ago a camera was a matrix. You can now write down where a robot is, what its sensor measured, and how a learned function is fitted to fill a gap the geometry leaves.

- **Coming with the papers that use them:** attention and transformers, networks that read point clouds directly, and sparse 3D convolution.
- **Next:** the representation catalog, point clouds, voxels and octrees, TSDF and ESDF, meshes, implicit surfaces, Gaussians, organised by what the robot has to *do* with them.
- **Lab 1 is out Fri Sep 18. Both diagnostic quizzes** are due **Sun Sep 27.**